

TRIGA Mark I

Independent neutronic and thermal-hydraulic
reference calculations in GeN-Foam and Cardinal

A code-to-code study of a 250 kW pool-type research reactor, comparing a deterministic OpenFOAM multiphysics solution (pcL) against a stochastic Monte Carlo reference (pcM).

pcL	GeN-Foam / OFFBEAT — 2-group diffusion, finite volume
pcM	Cardinal / OpenMC / MOOSE / NekRS — continuous-energy Monte Carlo
Nuclear data	ENDF/B-VIII.0 (both); ENDF/B-VII.1 for the pcM coupled run
Repository	github.com/kavyawadhwa134/triga-markI-benchmark
Independence	pcM frozen at tag <code>pcm-independent-v1</code> before pcL was read

STATUS OF THIS DOCUMENT

This report presents **verification and a cross-code comparison**. It is **not a validation report**. No experimental data was used in either calculation, and agreement between two codes is not a substitute for measurement. Section 5 states the status of each verification level explicitly.

Contents

1	Problem definition	2
1.1	Objective	2
1.2	Specification	2
1.3	The two approaches	2
1.4	Acceptance criterion	3
2	Results: pcL — GeN-Foam / OFFBEAT	4
2.1	Solver acceptance	4
2.2	Lattice constants	4
2.3	Steady state	4
2.4	Reactivity insertion transient	6
3	Results: pcM — Cardinal / OpenMC / MOOSE / NekRS	7
3.1	Method	7
3.2	Neutronics	7
3.3	Fuel temperature coefficient	7
3.4	Coupled neutronics–conduction	8
3.5	Thermal-hydraulics	9
4	Comparison	11
4.1	Overview	11
4.2	Case 1 — the directly comparable quantity	11
4.3	Cases 2–4 — the k_{eff} difference is method, not error	11
4.4	Case 5 — fuel temperature coefficient	12
4.5	Cases 6–8 — thermal quantities	13
4.6	Cases 10–12 — no pcM counterpart	13
5	Verification status	14
6	Conclusions	14

1 Problem definition

1.1 Objective

Two independent calculations of the same TRIGA Mark I configuration were performed using fundamentally different numerical methods, with the goal of establishing a cross-code reference. Neither team had access to the other’s implementation or results during the calculation; only the scientific problem definition was shared.

The TRIGA Mark I is a pool-type research reactor whose defining characteristic is uranium–zirconium hydride fuel. Hydrogen bound in the ZrH lattice produces a large, prompt, negative fuel temperature coefficient, which is the physical basis of the reactor’s inherent safety. Reproducing that feedback correctly is the principal physics requirement on any TRIGA model.

1.2 Specification

Table 1. Shared problem specification. Values marked *assumed* were not given and are documented modelling choices.

Parameter	Value	Basis
Fuel outer radius	1.815 cm	specification
Clad outer radius	1.866 cm	specification
Lattice pitch (square)	4.235 cm	specification
Active height	60.0 cm	specification
Core	0.495 m dia. $\times$ 0.60 m	single zone, vacuum boundaries
Fuel U–ZrH _{1.65}	6.0 g cm ^{−3}	U-235 1.7, U-238 6.8, Zr 89.862, H-1 1.638 wt%
Clad SS304	8.00 g cm ^{−3}	density <i>assumed</i>
Coolant	1.0 g cm ^{−3}	H ₂ O, unborated
Thermal power	250 kW	specification
Coolant inlet	293 K, 0.2 m s ^{−1}	specification
Pressure	0.1 MPa	<i>assumed</i> — open pool
Energy groups	2, cutoff 0.625 eV	specification

There is no fuel–clad gap: the specification gives two radii, so the clad inner radius equals the fuel outer radius. Thermal scattering $S(\alpha, \beta)$ data for `c_H_in_ZrH` and `c_Zr_in_ZrH` is mandatory; without it the bound-hydrogen physics that defines a TRIGA is absent.

1.3 The two approaches

Table 2. The two calculations solve different equations. This distinction governs the interpretation of every comparison in Section 4.

	pcL (GeN-Foam / OFFBEAT)	pcM (Cardinal)
Neutron transport	Deterministic, 2-group diffusion	Monte Carlo, continuous energy
Energy treatment	2 groups, 0.625 eV cutoff	Continuous, full spectrum
Core geometry	Homogenised single zone	Heterogeneous (+ homogenised variant)
Thermal–fluid	Finite volume, OpenFOAM	Spectral element, implicit LES
Fuel performance	OFFBEAT	Not modelled
Transient	Yes (+0.5\$ insertion)	Out of scope
Solution error	Discretisation, not quantified	Statistical, quantified ($\pm$ pcm)

1.4 Acceptance criterion

The project sets **500 pcm on k** as the flag for a material difference. Section 4 demonstrates this criterion is only meaningful between methodologically comparable quantities.

2 Results: pcL — GeN-Foam / OFFBEAT

2.1 Solver acceptance

Prior to the TRIGA calculation, GeN-Foam was verified against its own reference cases.

Table 3. pcL solver acceptance (`validation_step2.csv`).

Case	Quantity	Simulated	Expected	Status
3D SmallESFR	k_{eff}	0.936873	0.936827	PASS (4.9×10^{-5})
2D point kinetics	steady state	—	—	PASS
2D point kinetics	transient, no driveline	—	—	PASS
2D point kinetics	transient, driveline	—	—	PASS
2D point kinetics	transient, boron	—	—	PASS

2.2 Lattice constants

pcL condensed its own two-group constants from an OpenMC pin-cell calculation (70 batches, 20 inactive, 8000 particles per batch), giving $k_{\infty} = 1.38810 \pm 0.00123$.

Table 4. pcL two-group lattice constants (`validation_step5.csv`).

Quantity	Fast	Thermal
Removal (m^{-1})	4.77341	8.49190
$\nu\Sigma_f$ (m^{-1})	0.379899	12.5219
Diffusion coefficient (m)	0.00749023	0.00158363
χ	1.0	0.0
Inverse velocity (s m^{-1})	6.62997×10^{-6}	3.31676×10^{-4}
$S_{1 \rightarrow 1}$	39.7524	
$S_{1 \rightarrow 2}$	4.22693	
$S_{2 \rightarrow 1}$	0.023426	
$S_{2 \rightarrow 2}$		201.932

INTERNAL METHOD BENCHMARK



The same file reports a GeN-Foam diffusion solution for the *identical* pin cell: `genfoam_smoke_keff` = 1.331, against the Monte Carlo value of 1.38810 — a difference of **−4114 pcm**. This is pcL’s own measurement of deterministic-versus-stochastic disagreement on the simplest available geometry, and it is *eight times* the 500 pcm acceptance criterion before any core-geometry effect enters.

2.3 Steady state

The core is modelled as a bare homogenised box with vacuum boundaries.

Table 5. pcL steady-state results (`validation_step6.csv`).

Quantity	Value	Note
k_{eff}	1.240544	eigenvalue diffusion, bare homogenised box, vacuum BCs
Power	250 kW	target
Max coolant temperature	306.47 K	fluid T
Max average fuel temperature	594.2 K	sub-scale pin average; well below the 1023 K limit
Power peaking (max/avg)	2.45	single zone; cosine-like, flat by construction

**Figure 1.** pcL GeN-Foam computational mesh (left) and steady-state power density (right).**Figure 2.** pcL steady-state temperature field. The maximum lies in the upper portion of the core, consistent with coolant heating as it rises through the channel.

2.4 Reactivity insertion transient

A $+0.5\%$ ramp ($0.00325 \Delta k/k$) was inserted between $t = 100$ and $t = 100.5$ s.




Figure 3. pcL transient response. Power peaks at 489.7 kW and settles to a Doppler-limited equilibrium of 304.9 kW; average fuel temperature rises 32.89 K.

Table 6. pcL transient summary.

Quantity	Value	Note
Peak power	489 723 W	at $t = 100.56$ s
Final power	304 884 W	Doppler-limited equilibrium
Final average fuel T	544.59 K	at $t = 106$ s
Net reactivity	-3.87 pcm	external $+325$ vs Doppler -328.9

FEEDBACK COEFFICIENT IS AN INPUT, NOT A RESULT

The stated Doppler feedback of -328.9 pcm accompanies a fuel temperature rise from 511.70 to 544.59 K, a span of 32.89 K:

$$\frac{-328.9 \text{ pcm}}{32.89 \text{ K}} = -10.000 \text{ pcm/K.}$$

The result is *exactly* -10.000 , which indicates a constant coefficient supplied as a model input rather than one computed from nuclear data. Its origin is not documented in the committed files. Section 4 compares it against pcM's computed value.

3 Results: pcM — Cardinal / OpenMC / MOOSE / NekRS

3.1 Method

Continuous-energy Monte Carlo neutronics (OpenMC) is coupled to finite-element heat conduction (MOOSE) through Cardinal by Picard iteration: cell-averaged fission power is transferred to the solid, and cell temperatures are returned to the neutronics. The coolant is treated separately by spectral-element large-eddy simulation (NekRS).

Because the available hardware could not support a full-core coupled model, the coupled calculation uses a **single-pin unit cell**. All coupled temperatures therefore describe an *average* pin.

3.2 Neutronics

Table 7. pcM eigenvalue results, ENDF/B-VIII.0, continuous-energy Monte Carlo.

Quantity	Model	Value	Uncertainty
k_{∞}	2-D unit cell, reflective	1.384207	± 21 pcm
k_{eff}	3-D bare core, heterogeneous	1.102059	± 26 pcm
k_{eff}	3-D bare core, homogenised	1.145259	± 74 pcm
Leakage fraction	bare core	0.20149	± 0.00012
Power peaking (max/avg)	120 pins	2.2544	—
Peak power density	bare core	6.595 W cm^{-3}	—

The homogenised variant was computed specifically to separate geometry treatment from solution method in the comparison of Section 4. Its sensitivity to the $S(\alpha, \beta)$ choice is itself notable: 1.145259 (H in H_2O), 1.132505 (H in ZrH) and 1.136651 (none) span 1275 pcm, more than twice the acceptance criterion.

3.3 Fuel temperature coefficient

The defining TRIGA parameter was computed from first principles by sweeping fuel temperature with the moderator held at 293 K, isolating fuel feedback.




Figure 4. pcM fuel temperature coefficient. k_{∞} falls from 1.384926 at 296 K to 1.304249 at 1000 K.

Table 8. pcM local fuel temperature coefficient by interval.

Band	Coefficient
296–400 K	−5.08 pcm/K
400–500 K	−6.24 pcm/K
500–600 K	−7.76 pcm/K
600–800 K	−6.78 pcm/K
800–1 000 K	−5.92 pcm/K
296–1 000 K average	−6.34 pcm/K

The coefficient is large, negative, and becomes *more* negative as the fuel heats. Both features identify this as hydride physics rather than U-238 Doppler: a UO_2 LWR pin gives roughly -2 to -3 pcm/K and its magnitude does not grow in this way. The highlighted band is the one in which pcL’s transient operates.

3.4 Coupled neutronics–conduction




Figure 5. pcM coupled temperature solution. Left: axial profile, slightly top-skewed (305.98 K outlet against 302.56 K inlet) because the coolant heats as it rises. Right: radial profile at mid-plane; temperature is continuous across the fuel–clad interface at 1.815 cm, confirming the stitched contact.

Table 9. pcM coupled results and temperature drop budget.

Quantity	Value	Analytic	Share
Peak fuel temperature	359.0 K	362.8 K	—
Peak clad temperature	336.2 K	—	—
Power integral	2 083.33 W	2 083.33 W	exact
Film drop (clad → coolant)	40.3 K	42.4 K	63%
Fuel conduction drop	22.4 K	24.1 K	35%
Clad conduction drop	1.07 K	1.5 K	2%

Agreement with the analytic cylindrical-conduction solution is within 1.1%; the residual is explained by the true axial power shape being flatter than the cosine the hand calculation assumes. The film drop dominates and is the least certain term — it derives from a Dittus–Boelter correlation at $\text{Re} = 4755$, which is transitional and below the range where that correlation is reliable.




Figure 6. pcM Picard convergence. k_{eff} settles within Monte Carlo noise after the first temperature feedback; average fuel temperature is stable to 10^{-8} K.

3.5 Thermal-hydraulics

An implicit large-eddy simulation of the coolant subchannel was run at $\text{Re} = 4755$, polynomial order 7, 2880 spectral elements, on an NVIDIA GTX 1080 Ti. The run completed 20 000 steps but covered only **1.17 flow-through times**.




Figure 7. pcM NekRS evolution. Upper: coolant temperatures, all still rising at the final step, with the minimum below the 293 K inlet throughout. Lower: energy closure reaches only 34.6% of the steady balance.

Table 10. pcM thermal-hydraulic state at 20 000 steps.

Quantity	Value	Steady target	Status
Bulk outlet temperature	294.24 K	296.57 K	lower bound
Maximum temperature	333.03 K	—	still rising
Minimum temperature	291.52 K	≥ 293 K	unphysical
Advected heat / wall heat	34.6%	100%	not closed
Mass-flow imbalance	0.122%	$\rightarrow 0$	converged

THE PCM THERMAL-HYDRAULIC FIELD IS NOT CONVERGED

The flow field has converged — mass imbalance is flat at 0.122% from step 1400 onward — but only 34.6% of the wall heat leaves the domain, so 65.4% is still accumulating. **The coolant temperatures above are lower bounds, not steady-state results.** Reaching approximately five flow-throughs would require about 85 700 steps and 5.1 h of GPU time.

A second defect is recorded rather than smoothed over: the minimum temperature lies below the inlet in a domain whose only thermal boundary conditions are a 293 K inlet and a heated wall. With no heat sink this cannot be physical. It is an undershoot from the high-pass filter stabilisation and it does not decay.

4 Comparison

4.1 Overview

Table 11. Case-by-case comparison. Only cases marked *yes* support a direct numerical difference; the acceptance criterion applies to those alone.

#	Case	pcL	pcM	Difference	Comparable
1	k_{∞} , unit cell	1.388100	1.384207	−280 pcm	yes
2	k_{eff} , heterogeneous	1.240544	1.102059	+12566 pcm	no — method
3	k_{eff} , homogenised	1.240544	1.145259	+8320 pcm	partial
4	Diffusion vs. Monte Carlo	1.331	1.388100	−4114 pcm	pcL-internal
5	Fuel temp. coefficient	−10.000	−7.76	−29%	partial
6	Power peaking	2.45	2.2544	−8.0%	no
7	Max fuel temperature	594.2 K	359.0 K	−235 K	no
8	Max coolant temperature	306.5 K	294.2 K	—	no
9	Total power	250 kW	250 kW	0	yes
10	Transient peak power	489.7 kW	not run	—	no
11	Two-group MGXS	full set	not generated	—	no
12	Solver acceptance	PASS	not recorded	—	no

4.2 Case 1 — the directly comparable quantity

Both sides solved the same 2-D reflective pin cell with OpenMC continuous-energy Monte Carlo and ENDF/B-VIII.0. This is the only case in which the two calculations solve the same equations on the same geometry.




Figure 8. k_{∞} comparison. Bars are $\pm 1\sigma$ counting statistics; the shaded band is the 500 pcm acceptance criterion centred on pcL.

The difference is **−280 pcm** at 3.1σ — statistically resolved but comfortably inside the acceptance criterion. pcM’s uncertainty is approximately six times tighter, reflecting a 25-fold larger particle budget.

Interpretation. Two independent implementations of the same specification, in the same code, agree. This is the strongest evidence in the study that the shared problem definition was interpreted consistently on both sides.

4.3 Cases 2–4 — the k_{eff} difference is method, not error

Both cores are bare boxes with vacuum boundaries: pcL’s `validation_step6.csv` records the model verbatim as “*eigenvalue diffusion, bare homogenized box vacuum BCs*”. **Geometry is**

therefore not the explanation. An earlier hypothesis that pcL had modelled a graphite reflector is contradicted by this metadata and is withdrawn.




Figure 9. Decomposition of the k_{eff} difference into a homogenisation step (Monte Carlo on both sides) and a solution-method step (same homogenised bare box).

Table 12. Two-step decomposition of the core eigenvalue difference.

Step	From → to	$\Delta\rho$	What it isolates
1. Homogenisation	pcM het. → pcM homog.	+3920 pcm	Monte Carlo both sides; pure geometry smearing
2. Solution method	pcM homog. → pcL	+8320 pcm	Same homogenised bare box; pure method
Total	pcM het. → pcL	+12566 pcm	—

The residual +8320 pcm is corroborated independently by pcL’s own smoke test (Case 4), which measures −4114 pcm between GeN-Foam diffusion and Monte Carlo on the pin cell. The sign reverses between the two geometries — diffusion under-predicts on the reflective pin cell but over-predicts on the bare core — which is consistent with leakage treatment dominating in the bare case. pcM’s core leaks 20.1% of its neutrons, a regime in which few-group diffusion theory is least reliable.

A further uncontrolled variable applies to Cases 2–3: the project workflow intended pcM to generate the two-group MGXS and supply it to pcL, but **that phase was not run**. pcL condensed its own constants independently (Case 11). Re-solving pcL’s diffusion model with pcM-generated cross sections would isolate the remaining method error cleanly.

THE ACCEPTANCE CRITERION DOES NOT APPLY ACROSS METHODS

A 500 pcm threshold cannot meaningfully be applied between a continuous-energy Monte Carlo eigenvalue and a two-group diffusion eigenvalue. They are different quantities, and pcL’s own data places the method difference at thousands of pcm on even the simplest geometry. Applying the criterion across methods would manufacture a discrepancy that is not evidence of error in either calculation.

4.4 Case 5 — fuel temperature coefficient

This is the most consequential physics difference, because the coefficient governs the entire transient response.

Table 13. Fuel temperature coefficient, compared in pcL’s own operating band.

	Value	Origin
pcM, 500–600 K	−7.76 pcm/K	computed from $S(\alpha, \beta)$ ZrH data
pcM, 296–1 000 K average	−6.34 pcm/K	computed
pcL, transient	−10.000 pcm/K	constant input; origin undocumented
Difference in pcL’s band	−29%	pcL more negative

RECOMMENDED ACTION

A coefficient 29% more negative shuts a power excursion down faster. pcL’s peak power of 489.7 kW is therefore plausibly **under-predicted**, and its equilibrium reached sooner than a computed coefficient would give. For a TRIGA, where the prompt negative coefficient *is* the safety case, this warrants resolution.

pcM’s value is traceable to a nuclear-data calculation; pcL’s is not traceable from the committed files. pcL should either substitute a temperature-dependent coefficient and re-run the transient, or document the provenance of −10 pcm/K.

4.5 Cases 6–8 — thermal quantities

These cannot be differenced as reported, for three independent reasons:

1. **pcM describes an average pin.** Its own peaking factor of 2.2544 implies a hot channel is materially hotter, so the two fuel temperatures do not describe the same object.
2. **pcM’s film coefficient is a correlation, not a solution.** Dittus–Boelter at $Re = 4755$ supplies 63% of pcM’s total temperature rise, and is applied outside its reliable range.
3. **pcM’s coolant field is not converged** (34.6% energy closure), so 294.2 K is a lower bound rather than a result.

4.6 Cases 10–12 — no pcM counterpart

The transient is out of pcM’s defined scope: the project workflow specifies a steady-state production run and defines no transient phase. This is a documented scope decision rather than an omission, but it means pcL’s transient results — including the peak power on which the safety argument rests — have **no independent check**.

5 Verification status

THIS WORK IS VERIFIED, NOT VALIDATED

Verification asks whether the equations are solved correctly; validation asks whether the equations describe reality, and is establishable only against experimental data. **No experimental data was used in either calculation**, so no part of this study is validated. Agreement between two codes is not a substitute — two codes can agree and both be wrong, particularly when they share a specification, a nuclear data library, or a simplification. Here both share a bare-core model with no reflector and no control rods, so **neither result describes a physical TRIGA at criticality**.

Table 14. Verification status by level.

Level	Status	Evidence
1. Software	MET	Both codes run to completion; pcL passes its reference cases; pcM conserves power exactly at 2083.33 W
2. Numerical	PARTIAL	pcM neutronics and conduction converged; pcM CFD thermal field is not; no mesh or order refinement study either side
3. Physics	PARTIAL	pcM conduction within 1.1% of analytic; ZrH feedback confirmed; but pcM CFD shows an unphysical sub-inlet temperature and pcL's feedback coefficient is an undocumented input
4. Cross-code	PARTIAL	Agreement to 280 pcm where methodologically comparable; larger differences decomposed and attributed to homogenisation and method
5. Experimental	NOT ATTEMPTED	No experimental data exists in this project

Table 15. Actions required to close each gap, in priority order.

Gap	Action	Cost
1 pcL feedback coefficient	Document or recompute the assumed -10 pcm/K	—
2 k_{eff} method gap	Generate pcM's two-group MGXS; resolve pcL's diffusion with it	2 h CPU
3 pcM CFD not converged	Extend to ~ 5 flow-throughs after resolving the filter undershoot	5 h GPU
4 pcM sub-inlet temperature	Reduce scalar HPFRT strength; verify with a short run	1 h GPU
5 No mesh independence	Repeat at orders 5 and 8; compare integral quantities	8 h GPU
6 pcM library deviation	Re-run coupled production on ENDF/B-VIII.0	1 h CPU
7 Reflector and rods	Agree whether they enter the benchmark	—
8 No validation	Obtain TRIGA experimental data	—

6 Conclusions

1. **Where the two calculations are methodologically comparable, they agree.** The Monte Carlo pin cell differs by 280 pcm at 3.1σ , inside the 500 pcm acceptance criterion.

This is the meaningful cross-code result of the study.

2. **The core eigenvalue difference is a method difference, not an error.** It decomposes into 3920 pcm of homogenisation and 8320 pcm of diffusion-versus-Monte-Carlo on geometrically identical bare cores. pcL's own data independently measures a 4114 pcm method difference on the pin cell. Neither calculation is thereby shown to be incorrect.
3. **The acceptance criterion is not transferable across methods.** A 500 pcm threshold between a Monte Carlo and a diffusion eigenvalue would manufacture a false discrepancy.
4. **The fuel temperature coefficient is the outstanding physics question.** pcM computes -7.76 pcm/K in pcL's operating band from ZrH scattering data; pcL uses a constant -10.000 pcm/K of undocumented origin. Since this parameter governs the transient, and the transient has no independent check, it is the highest-priority item.
5. **Two results are incomplete and are reported as such.** pcM's thermal-hydraulic field has not converged, and pcM did not attempt the transient. Neither is presented as a result.
6. **Neither calculation is validated.** Both cores are bare, without reflector or control rods, and no experimental data was used. The study establishes cross-code consistency where the methods permit it — not fidelity to a physical reactor.

Data and inputs for both calculations: github.com/kavyawadhwa134/triga-markI-benchmark (branch `master`). pcM's full history and independence tag: github.com/kavyawadhwa134/pcm-triga-cardinal. pcM results were frozen at tag `pcm-independent-v1` before any pcL file was read.